

Big Sales Prediction using Random Forest Regressor

Get Understanding about Data set

There are 12 variables in dataset.

1. Item_Identifier
2. Item_Weight
3. Item_Fat_Content
4. Item_Visibility
5. Item_Type
6. Item_MRP
7. Outlet_Identifier
8. Outlet_Establishment_Year
9. Outlet_Size
10. Outlet_Location_Type
11. Outlet_Type
12. Item_Outlet_Sales

Import Library

```
In [1]: ▶ import pandas as pd
```

```
In [2]: ▶ import numpy as np
```

Import CSV as DataFrame

Use URL of file directly

```
In [3]: ▶ df = pd.read_csv('https://github.com/YBIFoundation/Dataset/raw/main/Big
```

Get the First Five Rows of Dataframe

```
In [4]: df.head()
```

```
Out[4]:
```

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP
0	FDT36	12.3	Low Fat	0.111448	Baking Goods	33.4874
1	FDT36	12.3	Low Fat	0.111904	Baking Goods	33.9874
2	FDT36	12.3	LF	0.111728	Baking Goods	33.9874
3	FDT36	12.3	Low Fat	0.000000	Baking Goods	34.3874
4	FDP12	9.8	Regular	0.045523	Baking Goods	35.0874

Get Information of DataFrame

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14204 entries, 0 to 14203
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Item_Identifier                       14204 non-null  object
1   Item_Weight                           11815 non-null  float64
2   Item_Fat_Content                       14204 non-null  object
3   Item_Visibility                       14204 non-null  float64
4   Item_Type                             14204 non-null  object
5   Item_MRP                              14204 non-null  float64
6   Outlet_Identifier                     14204 non-null  object
7   Outlet_Establishment_Year             14204 non-null  int64
8   Outlet_Size                           14204 non-null  object
9   Outlet_Location_Type                  14204 non-null  object
10  Outlet_Type                           14204 non-null  object
11  Item_Outlet_Sales                     14204 non-null  float64
dtypes: float64(4), int64(1), object(7)
memory usage: 1.3+ MB
```

Get Column Names

```
In [6]: df.columns
```

```
Out[6]: Index(['Item_Identifier', 'Item_Weight', 'Item_Fat_Content', 'Item_Visibility',
               'Item_Type', 'Item_MRP', 'Outlet_Identifier',
               'Outlet_Establishment_Year', 'Outlet_Size', 'Outlet_Location_Type',
               'Outlet_Type', 'Item_Outlet_Sales'],
              dtype='object')
```

Get the Summary Statistics

```
In [7]: df.describe()
```

```
Out[7]:
```

	Item_Weight	Item_Visibility	Item_MRP	Outlet_Establishment_Year	Item_Outlet_Sales
count	11815.000000	14204.000000	14204.000000	14204.000000	14204.000000
mean	12.788355	0.065953	141.004977	1997.830681	2185.000000
std	4.654126	0.051459	62.086938	8.371664	1827.000000
min	4.555000	0.000000	31.290000	1985.000000	31.290000
25%	8.710000	0.027036	94.012000	1987.000000	92.000000
50%	12.500000	0.054021	142.247000	1999.000000	176.000000
75%	16.750000	0.094037	185.855600	2004.000000	298.000000
max	30.000000	0.328391	266.888400	2009.000000	3122.000000

Get Missing Values Complete

```
In [8]: df['Item_Weight'].fillna(df.groupby(['Item_Type'])['Item_Weight'].transform('median'), inplace=True)
```

/tmp/ipykernel_18/3547494075.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['Item_Weight'].fillna(df.groupby(['Item_Type'])['Item_Weight'].transform('median'), inplace=True)
```

```
In [9]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14204 entries, 0 to 14203
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Item_Identifier                       14204 non-null  object
1   Item_Weight                           14204 non-null  float64
2   Item_Fat_Content                      14204 non-null  object
3   Item_Visibility                       14204 non-null  float64
4   Item_Type                             14204 non-null  object
5   Item_MRP                             14204 non-null  float64
6   Outlet_Identifier                     14204 non-null  object
7   Outlet_Establishment_Year            14204 non-null  int64
8   Outlet_Size                           14204 non-null  object
9   Outlet_Location_Type                 14204 non-null  object
10  Outlet_Type                           14204 non-null  object
11  Item_Outlet_Sales                     14204 non-null  float64
dtypes: float64(4), int64(1), object(7)
memory usage: 1.3+ MB
```

```
In [10]: df.describe()
```

```
Out[10]:
```

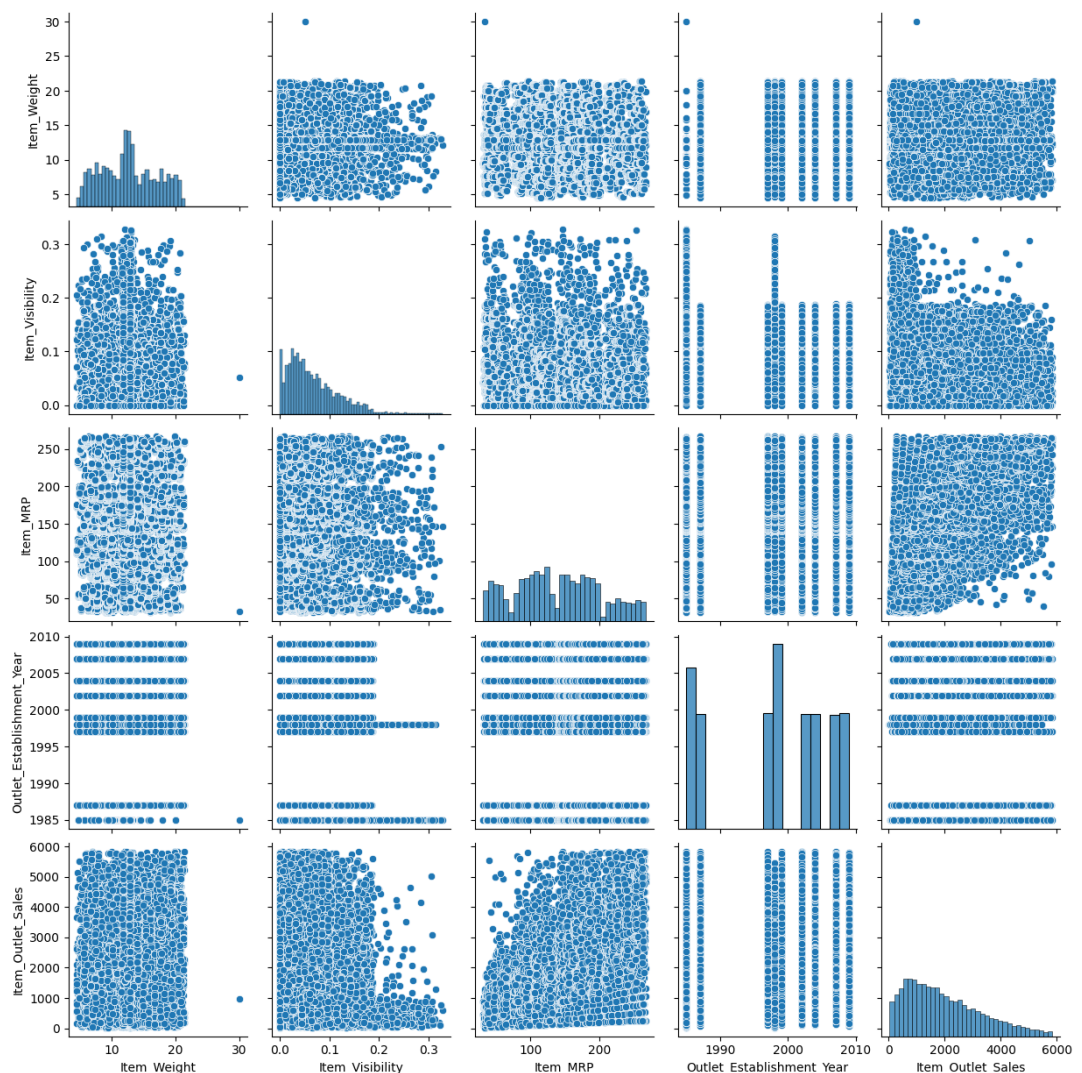
	Item_Weight	Item_Visibility	Item_MRP	Outlet_Establishment_Year	Item_Outlet_Sales
count	14204.000000	14204.000000	14204.000000	14204.000000	14204.000000
mean	12.742842	0.065953	141.004977	1997.830681	2185.000000
std	4.257583	0.051459	62.086938	8.371664	1827.000000
min	4.555000	0.000000	31.290000	1985.000000	31.290000
25%	9.300000	0.027036	94.012000	1987.000000	92.000000
50%	12.600000	0.054021	142.247000	1999.000000	176.000000
75%	16.000000	0.094037	185.855600	2004.000000	298.000000
max	30.000000	0.328391	266.888400	2009.000000	3122.000000

```
In [11]: # remove outlier
from scipy import stats
df = df[np.abs(stats.zscore(df['Item_Outlet_Sales'])) < 2]
```

```
In [12]: # pair plot
import seaborn as sns
sns.pairplot(df)
```

/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
 with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
 with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
 with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
 with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
 with pd.option_context('mode.use_inf_as_na', True):

Out[12]: <seaborn.axisgrid.PairGrid at 0x7bdec6aef4f0>



Get Categories and Counts of Categorical Variables

```
In [13]: df[['Item_Identifier']].value_counts()
```

```
Out[13]: Item_Identifier
FDQ08      10
FDD58      10
FDD56      10
FDD53      10
FDY36      10
..
FDC02       5
FDY55       5
NCL42       5
FDT21       5
FDA15       3
Name: count, Length: 1559, dtype: int64
```

```
In [14]: df[['Item_Fat_Content']].value_counts()
```

```
Out[14]: Item_Fat_Content
Low Fat      8173
Regular      4665
LF           512
reg          190
low fat      170
Name: count, dtype: int64
```

```
In [15]: df.replace({'Item_Fat_Content': {'LF': 'Low Fat', 'reg': 'Regular', 'low f
```

```
In [16]: df[['Item_Fat_Content']].value_counts()
```

```
Out[16]: Item_Fat_Content
Low Fat      8855
Regular      4855
Name: count, dtype: int64
```

```
In [17]: df.replace({'Item_Fat_Content': {'Low Fat': 0, 'Regular' : 1}}, inplace=
```

```
/tmp/ipykernel_18/722552587.py:1: FutureWarning: Downcasting behavior
in `replace` is deprecated and will be removed in a future version. To
retain the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to the future behavior, set `pd.set_option('future.no
_silent_downcasting', True)`
  df.replace({'Item_Fat_Content': {'Low Fat': 0, 'Regular' : 1}}, inplace=
```

```
In [18]: df[['Item_Type']].value_counts()
```

```
Out[18]: Item_Type
Fruits and Vegetables    1939
Snack Foods              1917
Household                1491
Frozen Foods             1372
Dairy                   1091
Baking Goods            1062
Canned                  1048
Health and Hygiene       828
Meat                    710
Soft Drinks              703
Breads                  406
Hard Drinks              347
Others                  273
Starchy Foods           261
Breakfast               177
Seafood                 85
Name: count, dtype: int64
```

```
In [19]: df.replace({'Item_Type':{'Fruits and Vegetables':0,'Snack Foods':0,'Hou
        'Frozen Foods' : 0, 'Dairy' : 0, 'Baking Goods
        'Canned' : 0, 'Health and Hygiene' : 1,
        'Meat' : 0, 'Soft Drinks' : 0, 'Breads' : 0, '
        'Others' : 2,'Starchy Foods' : 0, 'Breakfast'
        }},inplace=True)
```

/tmp/ipykernel_18/866940864.py:1: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future version. To retain the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`

```
df.replace({'Item_Type':{'Fruits and Vegetables':0,'Snack Foods':
0,'Household':1,
```

```
In [20]: df[['Item_Type']].value_counts()
```

```
Out[20]: Item_Type
0          11118
1           2319
2            273
Name: count, dtype: int64
```

```
In [21]: df[['Outlet_Identifier']].value_counts()
```

```
Out[21]: Outlet_Identifier
OUT018      1529
OUT046      1529
OUT013      1525
OUT045      1523
OUT049      1520
OUT035      1517
OUT017      1511
OUT027      1284
OUT010       925
OUT019       847
Name: count, dtype: int64
```

```
In [22]: df.replace({'Outlet_Identifier':{'OUT027': 0, 'OUT013': 1,
                                           'OUT049' : 2, 'OUT046' : 3, 'OUT035' : 4,
                                           'OUT045' : 5, 'OUT018' : 6,
                                           'OUT017' : 7, 'OUT010' : 8, 'OUT019' : 9,
                                           }}, inplace=True)
```

/tmp/ipykernel_18/2523128275.py:1: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future version. To retain the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`

```
df.replace({'Outlet_Identifier':{'OUT027': 0, 'OUT013': 1,
```

```
In [23]: df[['Outlet_Identifier']].value_counts()
```

```
Out[23]: Outlet_Identifier
3      1529
6      1529
1      1525
5      1523
2      1520
4      1517
7      1511
0      1284
8       925
9       847
Name: count, dtype: int64
```

```
In [24]: df[['Outlet_Size']].value_counts()
```

```
Out[24]: Outlet_Size
Medium      6768
Small       5417
High        1525
Name: count, dtype: int64
```



```
In [25]: df.replace({'Outlet_Size': {'Small': 0, 'Medium' : 1, 'High' : 1}}, inplace=True)
```

```
/tmp/ipykernel_18/171770719.py:1: FutureWarning: Downcasting behavior
in `replace` is deprecated and will be removed in a future version. To
retain the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to the future behavior, set `pd.set_option('future.no
_silent_downcasting', True)`
  df.replace({'Outlet_Size': {'Small': 0, 'Medium' : 1, 'High' : 1}}, i
nplace=True)
```

```
In [26]: df[['Outlet_Size']].value_counts()
```

```
Out[26]: Outlet_Size
1          8293
0          5417
Name: count, dtype: int64
```

```
In [27]: df[['Outlet_Location_Type']].value_counts()
```

```
Out[27]: Outlet_Location_Type
Tier 3          5263
Tier 2          4551
Tier 1          3896
Name: count, dtype: int64
```

```
In [28]: df.replace({'Outlet_Location_Type': {'Tier 1': 0, 'Tier 2' : 1, 'Tier 3' : 2}}, inplace=True)
```

```
/tmp/ipykernel_18/941750987.py:1: FutureWarning: Downcasting behavior
in `replace` is deprecated and will be removed in a future version. To
retain the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to the future behavior, set `pd.set_option('future.no
_silent_downcasting', True)`
  df.replace({'Outlet_Location_Type': {'Tier 1': 0, 'Tier 2' : 1, 'Tier
3' : 2}}, inplace=True)
```

```
In [29]: df[['Outlet_Location_Type']].value_counts()
```

```
Out[29]: Outlet_Location_Type
2          5263
1          4551
0          3896
Name: count, dtype: int64
```

```
In [30]: df[['Outlet_Type']].value_counts()
```

```
Out[30]: Outlet_Type
Supermarket Type1    9125
Grocery Store        1772
Supermarket Type2    1529
Supermarket Type3    1284
Name: count, dtype: int64
```

```
In [31]: df.replace({'Outlet_Type': {'Grocery Store': 0, 'Supermarket Type1' : 1,
```

/tmp/ipykernel_18/1904487509.py:1: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future version. To retain the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`

```
df.replace({'Outlet_Type': {'Grocery Store': 0, 'Supermarket Type1' : 1, 'Supermarket Type2' : 2, 'Supermarket Type3': 3}}, inplace=True)
```

```
In [32]: df[['Outlet_Type']].value_counts()
```

```
Out[32]: Outlet_Type
1          9125
0          1772
2          1529
3          1284
Name: count, dtype: int64
```

Define y (dependent or label or target variable) and X (independent or features or attribute Variable)

```
In [33]: # check correlation
df.describe().corr()
```

```
Out[33]:
```

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year	Outlet_Size	Outlet_Location_Type	Outlet_Type	Item_Outlet_Sales
Item_Weight	1.000000	0.999999	0.999999	0.999999	0.999907	0.999999	0.986922	0.999999	0.999999	0.999999	0.922230
Item_Fat_Content	0.999999	1.000000	1.000000	1.000000	0.999885	1.000000	0.986804	1.000000	1.000000	1.000000	0.921632
Item_Visibility	0.999999	1.000000	1.000000	1.000000	0.999884	1.000000	0.986805	1.000000	1.000000	1.000000	0.921613
Item_Type	0.999999	1.000000	1.000000	1.000000	0.999885	1.000000	0.986804	1.000000	1.000000	1.000000	0.921652
Item_MRP	0.999907	0.999885	0.999884	0.999885	1						
Outlet_Identifier	0.999999	1.000000	1.000000	1.000000							
Outlet_Establishment_Year	0.986922	0.986804	0.986805	0.986804							
Outlet_Size	0.999999	1.000000	1.000000	1.000000							
Outlet_Location_Type	0.999999	1.000000	1.000000	1.000000							
Outlet_Type	0.999999	1.000000	1.000000	1.000000							
Item_Outlet_Sales	0.922230	0.921632	0.921613	0.921652							

```
In [34]: y = df['Item_Outlet_Sales']
```

```
In [35]: X = df[['Item_Weight', 'Item_Fat_Content', 'Item_Visibility',  
               'Item_Type', 'Item_MRP', 'Outlet_Identifier',  
               'Outlet_Establishment_Year', 'Outlet_Size', 'Outlet_Location_Type',  
               'Outlet_Type']]
```

or use `.drop` function to define X

```
In [36]: X = df.drop(['Item_Identifier', 'Item_Outlet_Sales'], axis=1)
```

Get Train Test Split

```
In [37]: from sklearn.model_selection import train_test_split
```

```
In [38]: X_train, X_test, y_train, y_test = train_test_split(X,y, test_size = 0.2)
```

```
In [39]: X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
Out[39]: ((10968, 10), (2742, 10), (10968,), (2742,))
```

Get Model Train

```
In [40]: from sklearn.ensemble import RandomForestRegressor
```

```
In [41]: rfr = RandomForestRegressor()
```

```
In [42]: rfr.fit(X_train, y_train)
```

```
Out[42]: RandomForestRegressor()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

Get Model Prediction

```
In [43]: y_pred = rfr.predict(X_test)
```

Get Model Evaluation

```
In [44]:  from sklearn.metrics import mean_absolute_error
```

```
In [45]:  mean_absolute_error(y_test, y_pred)
```

```
Out[45]: 692.0422332454967
```

Get Visualization of Actual Vs Predicted Results

```
In [46]:  import matplotlib.pyplot as plt
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual Price vs Preicted Price")
plt.show()
```

